

Apache Spark Workshop

J.T. Halbert, Markus Dale, Jason Morris

February 24, 2015

Workshop goals

- ▶ Explore the Enron Dataset (<https://www.cs.cmu.edu/~./enron/>) while learning to use Apache Spark. Hopefully we can answer three questions:
 - ▶ Which ENRON employees spent too much time organizing their emails?
 - ▶ What hour of the week were most emails sent?
 - ▶ When did the FBI show up?
- ▶ Learn to appreciate the Scala Collections Library and the Resilient Distributed Dataset

Workshop non-goals

- ▶ Deploy and monitor Spark Applications
- ▶ Graphs
- ▶ MLlib

Overview of resources

- ▶ This presentation and “answers” to the “exercises” are available at <https://github.com/jt-halbert/spark-workshop/>.
- ▶ The scripts we used to create the AWS EMR instances is available at <https://github.com/notjasonmorris/AWS>.
- ▶ The ETL code that is making our exploration of the Enron dataset so convenient is available at <https://github.com/medale/spark-mail/>.
- ▶ LARGE PORTIONS of this presentation are pulled from Markus' work. Thanks Markus!!

Who are we?

- ▶ I am Tetra's Chief Data Scientist and I help certain people learn certain things about certain parts of their data.
- ▶ Tetra Concepts, LLC is the finest collection of development talent anyone could ask for (AND SOME OF THEM ARE WALKING AMONG YOU!!!!11!!1!!!!)

Why Apache Spark?

- ▶ Excellent question: four parts of a partial answer.

Data Science is a filthy job

- ▶ I am not even sure what Data Science is. They put Science right in the name, so it must be pretty serious right?
- ▶ I like to think it is the disciplined application of a scientific mindset to that nebulous thing called “data.”
- ▶ But really...
 - ▶ You spend 90% of your time getting and cleaning that thing.
 - ▶ And when you finally get it cleaned and available it very often is unwieldy in some further way (size or speed.)
 - ▶ The act of cleaning itself is a data science problem.

The Ecosystem is big and filled with snakes

- ▶ Berg Data

You can lie with statistics

- ▶ In a big enough database you can find a set of columns to perfectly predict any outcome.
 - ▶ <http://www.tylervigen.com/>
- ▶ We need better tools so we can spend more time doing the hard work of telling the truth (or some approximation).

Why Apache Spark?

- ▶ Spark gives you a way to explore small, medium, large, (very large ?) data in a convenient way.
 - ▶ You can actually explore distributed datasets: lazy evaluation and a rich collections api.
 - ▶ You can scale your exploratory code up to a full job relatively quickly: REPL driven development.
- ▶ It wraps an increasing amount of the Hadoop Ecosystem and plays naturally.

Your customer wants pretty little magical things



Figure 1: Spark wraps a lot of other peoples toys

Let's get started

- ▶ The first step is to learn enough Scala to be dangerous.

Combinator functions on Scala collections

- ▶ Examples: map, flatMap, filter, reduce, fold, aggregate
- ▶ Background - Combinatory logic, higher-order functions...

Combinatory Logic

Moses Schönfinkel and Haskell Curry in the 1920s

[C]ombinator is a higher-order function that uses only function application and earlier defined combinators to define a result from its arguments [Combinatory Logic @wikipedia_combinatory_2014]

A *Higher-Order Function* is a function that takes functions as arguments or returns function.

map

- ▶ Applies a given function to every element of a collection
- ▶ Returns collection of outputs of that function
- ▶ input argument - same type as collection type
- ▶ return type - can be any type

map - Scala

```
def computeLength(w: String): Int = w.length

val words = List("when", "shall", "we", "three",
  "meet", "again")
val lengths = words.map(computeLength)

> lengths    : List[Int] = List(4, 5, 2, 5, 4, 5)
```


map - Scala syntactic sugar

```
//anonymous function (specifying input arg type)  
val list2 = words.map((w: String) => w.length)
```

```
//let compiler infer arguments type  
val list3 = words.map(w => w.length)
```

```
//use positionally matched argument  
val list4 = words.map(_.length)
```

map - ScalaDoc

See immutable List ScalaDoc

List[+A]

...

```
final def map[B](f: (A) => B): List[B]
```

- ▶ Builds a new collection by applying a function to all elements of this list.
- ▶ B - the element type of the returned collection.
- ▶ f - the function to apply to each element.
- ▶ returns - a new list resulting from applying the given function f to each element of this list and collecting the results.

flatMap

- ▶ ScalaDoc:

```
List[+A]
```

```
...
```

```
def flatMap[B](f: (A) =>  
    GenTraversableOnce[B]): List[B]
```

- ▶ GenTraversableOnce - List, Array, Option...
- ▶ can be empty collection or None
- ▶ flatMap takes each element in the GenTraversableOnce and puts it in order to output List[B]
- ▶ removes inner nesting - flattens
- ▶ output list can be smaller or empty (if intermediates were empty)

flatMap Example

```
val macbeth = """When shall we three meet again?  
|In thunder, lightning, or in rain?""".stripMargin  
val macLines = macbeth.split("\n")  
// macLines: Array[String] = Array(  
//   When shall we three meet again?,  
//   In thunder, lightning, or in rain?)  
  
//Non-word character split  
val macWordsNested: Array[Array[String]] =  
    macLines.map{line => line.split("""\W+""")}  
//Array(Array(When, shall, we, three, meet, again),  
//      Array(In, thunder, lightning, or, in, rain))  
  
val macWords: Array[String] =  
    macLines.flatMap{line => line.split("""\W+""")}  
//Array(When, shall, we, three, meet, again, In,  
//      thunder, lightning, or, in, rain)
```

filter

```
List[+A]
```

```
...
```

```
def filter(p: (A) => Boolean): List[A]
```

- ▶ selects all elements of this list which satisfy a predicate.
- ▶ returns - a new list consisting of all elements of this list that satisfy the given predicate p. The order of the elements is preserved.

filter Example

```
val macWordsLower = macWords.map{_.toLowerCase}  
//Array(when, shall, we, three, meet, again, in, thunder,  
//      lightning, or, in, rain)  
  
val stopWords = List("in","it","let","no","or","the")  
val withoutStopWords =  
    macWordsLower.filter(word => !stopWords.contains(word))  
// Array(when, shall, we, three, meet, again, thunder,  
//      lightning, rain)
```

reduce

```
List[+A]
```

```
...
```

```
def reduce[A1 >: A](op: (A1, A1) => A1): A1
```

- ▶ Creates one cumulative value using the specified associative binary operator.
- ▶ A1 - A type parameter for the binary operator, a supertype (super or same) of A. (List is covariant +A)
- ▶ op - A binary operator that must be associative.
- ▶ returns - The result of applying op between all the elements if the list is nonempty. Result is same type as (or supertype of) list type.
- ▶ UnsupportedOperationException if this list is empty.

reduce Example

```
//beware of overflow if using default Int!  
val numberOfAttachments: List[Long] =  
    List(0, 3, 4, 1, 5)  
val totalAttachments =  
    numberOfAttachments.reduce((x, y) => x + y)  
//Order unspecified/non-deterministic, but one  
//execution could be:  
//0 + 3 = 3, 3 + 4 = 7,  
//7 + 1 = 8, 8 + 5 = 13  
  
val emptyList: List[Long] = Nil  
//UnsupportedOperationException  
emptyList.reduce((x, y) => x + y)
```


fold

List[+A]

...

```
def fold[A1 >: A](z: A1)(op: (A1, A1) => A1): A1
```

- ▶ Very similar to reduce but takes start value z (a neutral value, e.g. 0 for addition, 1 for multiplication, Nil for list concatenation)
- ▶ returns start value z for empty list
- ▶ Note: See also foldLeft/Right (return completely different type)

```
foldLeft[B](z: B)(f: (B, A) => B): B
```

fold Example

```
val numbers = List(1, 4, 5, 7, 8, 11)
val evenCount = numbers.fold(0) { (count, currVal) =>
  println(s"Count: $count, value: $currVal")
  if (currVal % 2 == 0) {
    count + 1
  } else {
    count
  }
}
```

Count: 0, value: 1

Count: 0, value: 4

Count: 1, value: 5

Count: 1, value: 7

Count: 1, value: 8

Count: 2, value: 11

evenCount: Int = 2

aggregate

```
List[+A]  
...  
def aggregate[B] (z: B) (seqop: (B, A) => B,  
                        combop: (B, B) => B): B
```

- ▶ More general than fold or reduce. Can return different result type.
- ▶ Apply seqop function to each partition of data.
- ▶ Then apply combop function to combine all the results of seqop.
- ▶ On a normal immutable list this is just a foldLeft with seqop (but on a parallelized list both operations are called).

aggregate Example

```
val wordsAll = List("when", "shall", "we", "three",  
    "meet", "again", "in", "thunder", "lightning",  
    "or", "in", "rain")  
//Map(5 letter words ->3, 9->1, 2->4, 7->1, 4->3)  
val lengthDistro = wordsAll.aggregate(Map[Int, Int]())(  
    seqop = (distMap, currWord) =>  
    {  
        val length = currWord.length()  
        val newCount = distMap.getOrElse(length, 0) + 1  
        val newKv = (length, newCount)  
        distMap + newKv  
    },  
    combop = (distMap1, distMap2) => {  
        distMap1 ++ distMap2.map {  
            case (k, v) =>  
                (k, v + distMap1.getOrElse(k, 0))  
        }  
    })
```

So what does this have to do with Apache Spark?

- ▶ Resilient Distributed Dataset (RDD)
- ▶ From API docs: “immutable, partitioned collection of elements that can be operated on in parallel”
- ▶ map, flatMap, filter, reduce, fold, aggregate...

Spark - RDD API

- ▶ RDD API
- ▶ Transforms - map, flatMap, filter, reduce, fold, aggregate...
 - ▶ Lazy evaluation (not evaluated until action!)
- ▶ Actions - count, collect, first, take, saveAsTextFile...

Spark - From RDD to PairRDDFunctions

- ▶ If an RDD contains tuples (K,V) - can apply PairRDDFunctions
- ▶ Uses implicit conversion of RDD to PairRDDFunctions
- ▶ In 1.2 and previous versions available by importing `org.apache.spark.SparkContext._`

From `org.apache.spark.SparkContext`:

```
implicit def rddToPairRDDFunctions[K, V](  
  rdd: RDD[(K, V)])  
  (implicit kt: ClassTag[K],  
   vt: ClassTag[V],  
   ord: Ordering[K] = null) = {  
    new PairRDDFunctions(rdd)  
  }
```

PairRDDFunctions

- ▶ keys, values - return RDD of keys/values
- ▶ mapValues - transform each value with a given function
- ▶ flatMapValues - flatMap each value (0, 1 or more output per value)
- ▶ groupByKey - `RDD[(K, Iterable[V])]`
 - ▶ Note: expensive for aggregation/sum - use `reduce/aggregateByKey`!
- ▶ reduceByKey - return same type as value type
- ▶ foldByKey - zero/neutral starting value
- ▶ aggregateByKey - can return different type
- ▶ join (left/rightOuterJoin), cogroup ...

From RDD to DoubleRDDFunctions

- ▶ From API docs: “Extra functions available on RDDs of Doubles through an implicit conversion. Import `org.apache.spark.SparkContext._`”

```
From org.apache.spark.SparkContext:  
implicit def doubleRDDToDoubleRDDFunctions(  
  rdd: RDD[Double])  
  = new DoubleRDDFunctions(rdd)
```

DoubleRDDFunctions

- ▶ mean, stddev, stats (count, mean, stddev, min, max)
- ▶ sum
- ▶ histogram

Analytic 1 - Mail Folder Statistics In MapReduce

- ▶ What are the least/most/average number of folders per user?
- ▶ Each MailRecord has user name and folder name

```
lay-k/      <- mailFields(UserName)
  business  <- mailFields(FolderName)
  family
  enron
  inbox
  ...
```

Hadoop Mail Folder Stats - Mapper

- ▶ read each mail record
- ▶ emits key: userName, value: folderName for each email

Hadoop Mail Folder Stats - Reducer

- ▶ reduce method
 - ▶ create set from values for a given key (unique folder names per user)
 - ▶ `set.size == folder count`
 - ▶ keep adding up all `set.size` (`totalNumberOfFolders`)
 - ▶ one up counter for each key (`totalUsers`)
 - ▶ keep track of min/max count
- ▶ cleanup method
 - ▶ compute average for this partition:
`totalNumberOfFolders/totalUsers`
 - ▶ write out min, max, `totalNumberOfFolders`, `totalUsers`, `avgPerPartition`

Hadoop Mail Folder Stats - Driver

- ▶ Set Input/OutputFormat
- ▶ Number of reducers

Hadoop Mail Folder Stats - Results

- ▶ if only one reducer - results are overall lowest/highest/avg
- ▶ if multiple reducers
 - ▶ post-processing overall lowest/highest
 - ▶ add totalNumberOfFolders and totalUsers to compute overall average

Hadoop Mapper

```
public void map(AvroKey<MailRecord> key,
NullWritable value, Context context) throws ... {
    MailRecord mailRecord = key.datum();
    Map<CharSequence, CharSequence> mailFields =
        mailRecord.getMailFields();
    CharSequence userName =
        mailFields.get(AvroMailMessageProcessor.USER_NAME);
    CharSequence folderName =
        mailFields.get(AvroMailMessageProcessor.FOLDER_NAME);
    userKey.set(userName.toString());
    folderValue.set(folderName.toString());
    context.write(userKey, folderValue);
}
```


Hadoop Reducer

```
public void reduce(Text userKey,
    Iterable<Text> folderValues,
    Context context) throws ... {
    Set<String> uniqueFolders = new HashSet<String>();
    for (Text folder : folderValues) {
        uniqueFolders.add(folder.toString());
    }
    int count = uniqueFolder.size();
    if (count > maxCount) maxCount = count;
    if (count < minCount) minCount = count;
    totalNumberOfFolder += count
    totalUsers++
}
...
public void cleanup...
//write min, max, totalNumberOfFolders,
//totalUsers, avgPerPartition
```

Let's get to work